

The objective of this study is to build a logistic regression model to predict the customer churn based on a set of 9 predictors. Refer below for code loading the data and establishing a random split into 70% training and 30% testing sample.

```
data = pd.read_csv("C:\\Users\\tgmce\\Downloads\\Bank Customer Churn Prediction_revised.csv")
x = data.iloc[:, :-1]
y = data.iloc[:, -1:]
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.30, random_state = 0)
```

After we import and split the data, we build a logistic regression model on the training data with all predictors.

```
model=sm.Logit(y_train, sm.add_constant(x_train))
lr = model.fit()
print(lr.summary())
```

```
Optimization terminated successfully.
      Current function value: 0.435680
      Iterations 6
```

```

                        Logit Regression Results
=====
Dep. Variable:          churn    No. Observations:          7000
Model:                  Logit    Df Residuals:              6990
Method:                  MLE     Df Model:                  9
Date:                    Wed, 02 Apr 2025    Pseudo R-squ.:          0.1348
Time:                    02:03:48    Log-Likelihood:         -3049.8
converged:               True     LL-Null:                 -3524.9
Covariance Type:         nonrobust    LLR p-value:            9.177e-199
=====

```

	coef	std err	z	P> z	[0.025	0.975]
const	-3.1303	0.288	-10.869	0.000	-3.695	-2.566
credit_score	-0.0009	0.000	-2.711	0.007	-0.002	-0.000
gender	-0.5389	0.065	-8.347	0.000	-0.665	-0.412
age	0.0715	0.003	23.475	0.000	0.066	0.077
tenure	-0.0283	0.011	-2.551	0.011	-0.050	-0.007
balance	5.012e-06	5.49e-07	9.130	0.000	3.94e-06	6.09e-06
products_number	-0.1165	0.057	-2.053	0.040	-0.228	-0.005
credit_card	-0.0186	0.070	-0.265	0.791	-0.156	0.119
active_member	-1.0000	0.068	-14.736	0.000	-1.133	-0.867
estimated_salary	6.968e-07	5.6e-07	1.244	0.214	-4.01e-07	1.8e-06

```
=====
```

```
model=sm.Logit(y_train, sm.add_constant(x_train))
lr = model.fit()
print(lr.summary())
```

As shown from the summary table above, the fitted linear regression function is:

- $\log(P(\text{churn}=1) / 1 - P(\text{churn}=1)) = -3.13 - 0.0009\text{credit_score} - 0.539\text{gender} + 0.072\text{age} - 0.028\text{tenure} + 5.012 \times 10^{-6}\text{balance} - 0.117\text{products_number} - 0.019\text{credit_card} - \text{active_member} + 6.97 \times 10^{-7}\text{estimated_salary}$

Two predictors are not significant given 0.05 significance level (See: credit_card and estimated_salary).

The coefficient of credit_score can be interpreted as for every one unit increase of customer score, the odds of churn will decrease by 0.09% ($=e^{-0.0009} - 1$). The coefficient of active_member can be interpreted as the odds of churn for active member is 63.21% ($=e^{-1} - 1$) lower than that for non-active member.

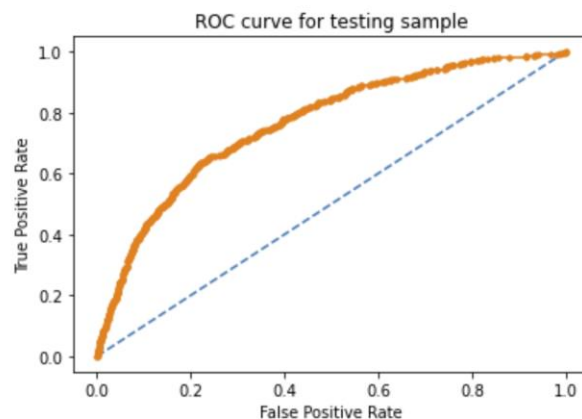
We can use the logistic regression model trained above to make prediction on both training and testing sample.

```
p_pred_train = lr.predict(sm.add_constant(x_train))
p_pred_test = lr.predict(sm.add_constant(x_test))
```

The ROC curve for the testing sample can be found below. AUC is 0.76 for the testing sample.

```
mr_probs = [0 for _ in range(len(y_test))]
mr_fpr, mr_tpr, _ = roc_curve(y_test, mr_probs)
lr_fpr, lr_tpr, _ = roc_curve(y_test, p_pred_test)
plt.plot(mr_fpr, mr_tpr, linestyle='--')
plt.plot(lr_fpr, lr_tpr, marker='.')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC curve for testing sample')
print('AUC for testing sample is:', metrics.auc(lr_fpr, lr_tpr))
```

AUC for testing sample is: 0.7648574246340937



The confusion matrix on the testing sample can be found below using 0.5 as the cutoff probability.

```

from sklearn.metrics import confusion_matrix
y_pred = round(p_pred_test)
print('Confusion matrix is:', confusion_matrix(y_test, y_pred))
tn, fp, fn, tp = confusion_matrix(y_test, y_pred).ravel()
print('MR:', (fp+fn)/(tn+fp+fn+tp))
print('TPR:', tp/(tp+fn))
print('FNR:', fn/(tp+fn))

```

```

Confusion matrix is: [[2305  74]
 [ 517 104]]
MR: 0.197
TPR: 0.16747181964573268
FNR: 0.8325281803542673

```

The misclassification rate is 19.7%, true positive rate is 16.75%, and false negative rate is 83.25%.

For feature selection I elected to use backward stepwise selection with the objective to minimize the value of AIC.

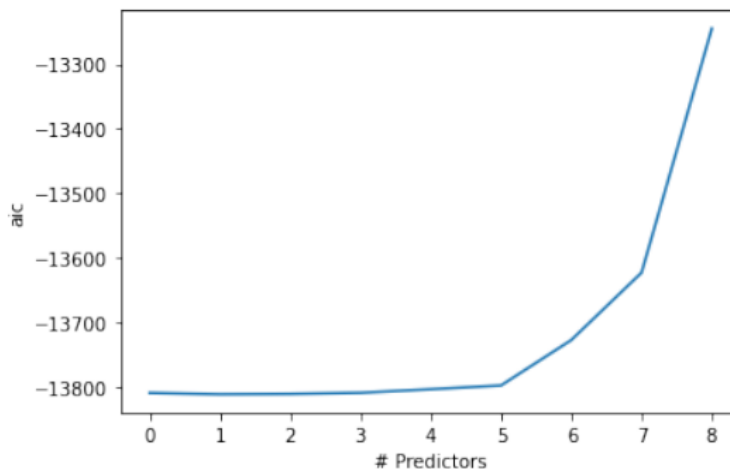
```

def fit_lr(feature_set):
    model = sm.Logit(y_train, sm.add_constant(x_train[list(feature_set)]))
    lr = model.fit()
    MSE = mean_squared_error(lr.predict(sm.add_constant(x_train[list(feature_set)])), y_train)
    AIC = len(x_train)*np.log(MSE)+2*len(lr.params)
    return {"model":lr, "MSE":MSE, "AIC":AIC}

def getbest(k):
    result = []
    for combo in itertools.combinations(x_train.columns, k):
        result.append(fit_lr(combo))
    models = pd.DataFrame(result)
    best_model = models.loc[models['MSE'].argmin()]
    return best_model

picked = list(x_train.columns)
models_best = pd.DataFrame(columns=["model", "MSE", "AIC"])
models_best.loc[0] = fit_lr(picked)
for i in range(1,9):
    best_MSE = np.inf
    for combo in itertools.combinations(picked, len(picked)-1):
        MSE = fit_lr(list(combo))["MSE"]
        if MSE < best_MSE:
            best_MSE = MSE
            best_feature = combo
    picked = best_feature
    models_best.loc[i] = fit_lr(list(picked))
print(models_best)
plt.plot(models_best["AIC"])
plt.xlabel('# Predictors')
plt.ylabel('aic')

```



As shown, the model with 1 predictors being removed (8 predictors included) is associated with the lowest AIC. All predictors but credit_card are included in the final model.

```
print(models_best.loc[1, "model"].summary())
```

```

Logit Regression Results
=====
Dep. Variable:          churn    No. Observations:          7000
Model:                  Logit    Df Residuals:              6991
Method:                  MLE     Df Model:                  8
Date:                   Wed, 02 Apr 2025    Pseudo R-squ.:          0.1348
Time:                   02:05:54    Log-Likelihood:         -3049.8
converged:              True     LL-Null:                -3524.9
Covariance Type:        nonrobust    LLR p-value:            8.441e-200
=====

```

	coef	std err	z	P> z	[0.025	0.975]
const	-3.1437	0.284	-11.087	0.000	-3.699	-2.588
credit_score	-0.0009	0.000	-2.712	0.007	-0.002	-0.000
gender	-0.5390	0.065	-8.350	0.000	-0.666	-0.413
age	0.0715	0.003	23.481	0.000	0.066	0.078
tenure	-0.0284	0.011	-2.558	0.011	-0.050	-0.007
balance	5.014e-06	5.49e-07	9.135	0.000	3.94e-06	6.09e-06
products_number	-0.1165	0.057	-2.052	0.040	-0.228	-0.005
active_member	-0.9996	0.068	-14.734	0.000	-1.133	-0.867
estimated_salary	6.97e-07	5.6e-07	1.244	0.214	-4.01e-07	1.8e-06

```
=====
```

We used the re-trained logistic regression to make prediction on the testing sample.

```

lr_backward = models_best.loc[1, "model"]
p_pred_test = lr_backward.predict(sm.add_constant(x_test[['credit_score', 'gender', 'age', 'tenure', 'balance', 'products_number',
mr_probs = [0 for _ in range(len(y_test))])
mr_fpr, mr_tpr, _ = roc_curve(y_test, mr_probs)
lr_fpr, lr_tpr, _ = roc_curve(y_test, p_pred_test)
print('AUC for testing sample is:', metrics.auc(lr_fpr, lr_tpr))

```

AUC for testing sample is: 0.7648574246340937

The AUC for the testing sample is 0.7649. My model has a 76.49% chance of ranking a randomly chosen positive case (e.g., customer did churn) higher than a randomly chosen negative case (e.g., customer did not churn).